

CPSC 457 T01/T04

Xining Chen

Deadlock Detector

5 -> 1

- Add 5 as a node into `std::unordered_map adj_list`
- Add 1 as a node into `std::unordered_map adj_list`
- Increment node 5 out count in `out_counts` list
- Missing a step?

Deadlock Detector

5 -> 1

- Add 5 as a node into `std::unordered_map adj_list`
- Add 1 as a node into `std::unordered_map adj_list`
- Increment node 5 out count in `out_counts` list
- `out_counts` is an `unordered_map`. Therefore node 1 is **not** recorded in `out_counts`.
- Need “dummy” / “filler” code – to add node 1 into `out_counts`, initialized with count to 0.

Deadlock Detector

- Q: Is it necessary to distinguish processes vs resources?
- A: Yes. When printing output, **only print processes**.
- Q: How do you distinguish them?
- A: Infinite number of ways. One option is:
 prepend the node name with string 'p_' and 'r_' for process and resources respectively. I.e. p_x is a process with name x, r_x is a process with name x.

p_x and r_x are clearly different strings. Hence hash values will be different.

Deadlock Detector


- Q: What's an edge index?
- A:

Few more examples:

```
1 $ cat test2.txt
```


```
0 5 -> 1  
1 5 <- 3  
2 5 -> 2  
3 7 <- 1  
4 12 <- 2  
5 7 -> 3 *
```

```
$ ./deadlock < test2.txt  
edge_index : 5  
cycle      : [5,7]
```



```
$ cat test3.txt
```

```
p7 <- r7 0  
p7 -> r1 1  
p3 -> r7 2  
p3 <- r1 3 *  
p12 <- r1 4  
$ ./deadlock < test3.txt  
edge_index : 3  
cycle      : [p7,p3]
```




Deadlock Detector

- Exercise: Draw and figure out the edge_index. Does it match the computed edge_index?

```
$ cat test4.txt
```

```
12 -> 1  
12 <- 1000  
7 -> 1000  
7 <- 1  
2 -> 3  
2 <- 432  
77 -> 432  
77 <- 3
```


```
$ ./deadlock < test4.txt  
edge_index : 3  
cycle      : [12,7]
```



```
$ cat test5.txt
```

```
12 -> 1  
12 <- 1000  
7 -> 1000  
2 -> 3  
2 <- 432  
77 -> 432  
77 <- 3  
7 <- 1
```

```
$ ./deadlock < test5.txt  
edge_index : 6  
cycle      : [2,77]
```



Simple “trivial” exercises are important. They are low time cost (1-3 mins) and ensures you truly do understand the concepts.

Deadlock Detector


- Q: In test5.txt, can I report the cycle 12, 7 instead of 2, 77? Would that be correct?
- A: No. That is incorrect because it is not the first deadlock. The **first deadlock** occurs on the right cycle. You must report the right cycle.

```
$ cat test5.txt
```

```
12 -> 1  
12 <- 1000  
7 -> 1000  
2 -> 3  
2 <- 432  
77 -> 432  
77 <- 3  
7 <- 1
```

```
$ ./deadlock < test5.txt
```


```
edge_index : 6  
cycle      : [2,77]
```



Deadlock Detector

- Question: Does the following process count as being stuck in a deadlock (& therefore should be printed?)

- A:



Yes! Process 81 is stuck in a deadlock.

```
$ cat myTest5.txt
```

```
12 -> 1
```

```
12 <- 100
```

```
7 -> 1000
```

```
2 -> 3
```

```
2 <- 432
```

```
77 -> 432
```

```
81 -> 432
```

```
77 <- 3
```

```
7 <- 1
```

```
$ ./deadlock < myTest5.txt
```

```
edge index : 6
```

```
cycle      : [2, 77, 81]
```


Deadlock Detector

- Q: My code is very slow on test6.txt. Why is that?
- A: Probably didn't implement optimization #1 for the topological sort. See tutorial slides CPSC 457 – Assignment 4.pdf.

Deadlock Detector

- Checks:

1. Before implementing topological sort, ensure your graph is processed correctly.

```
Result detect_deadlock(const std::vector<std::string> & edges);
```

```
class Graph {  
    std::unordered_map<std::string, std::vector<std::string>> adj_list;  
    std::unordered_map<std::string, int> out_counts;  
    ...  
} graph;
```

Are processes and nodes added to adj_list correctly?

Are the out_counts set correctly?

Can you traverse through both maps and print adj_list and out_counts?

Deadlock Detector

- Checks:

2. Implement topological sort. Run it on the **final** graph. Easier to debug if you have 1 topological sort per call, instead of 1 per edge added. (But edge_index will be wrong, but that's ok).

Goal here: ensure topological sort works.

```
out = out_counts # copy out_counts so that we can modify it
zeros[] = find all nodes in graph with outdegree == 0
```

```
while zeros is not empty:
```

```
    n = remove one entry from zeros[]
```

```
    for every n2 of adj_list[n]:
```

```
        out[n2] --
```

```
        if out[n2] == 0:
```

```
            append n2 to zeros[]
```

```
cycle[] = all nodes n that represent a process and out[n]>0
```

Deadlock Detector

- Checks:

3. Move `topological_sort()` call to a different line in your code. It should be called after an edge is processed and added to Graph.

Check that the `edge_index` output is correct.

Output of cycle is correct (**print only processes in cycle**)

Deadlock Detector

- Checks:
 4. Convert into integer graph.

```
class FastGraph {  
    std::vector<std::vector<int>> adj_list;  
    std::vector<int> out_counts;  
    ...  
} graph;
```

Everything else stays the same. Might have to change some data types of functions.

1 Extra step: convert from integer (hash value) back to string for cycle[].

Scheduler simulation

- Talk about Appendix 2.
- Switch to ipad.

Next time:

- Open tutorial for 1-1 help.